Introduction

Geometry of Data

August 25, 2020

CIFAR-10

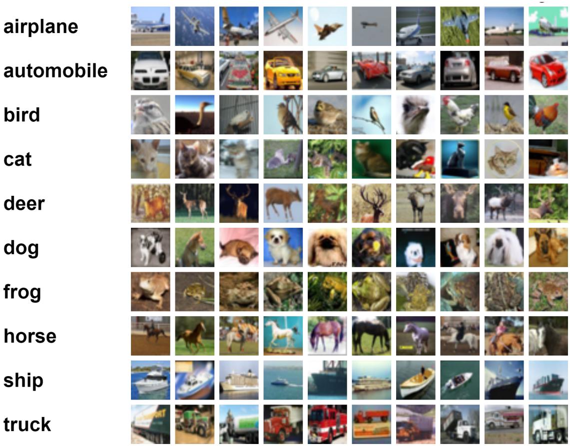

$32 \times 32 \times 3 = 3{,}072$ dimensions

10 classes

Uniform Random Images

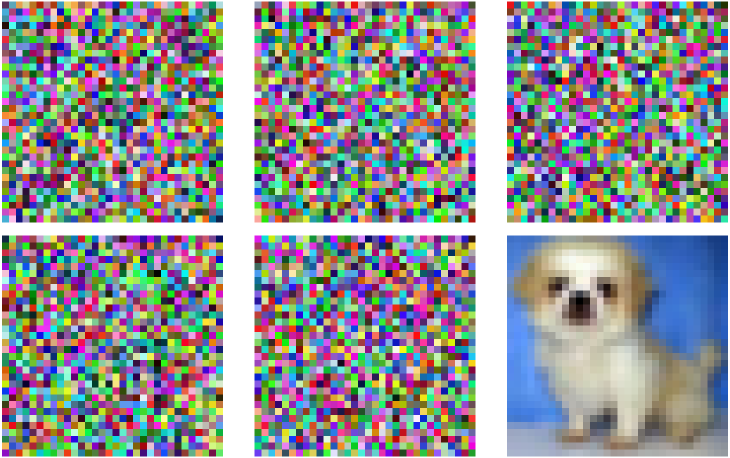

just kidding!

Manifold Hypothesis

Real data lie near lower-dimensional manifolds

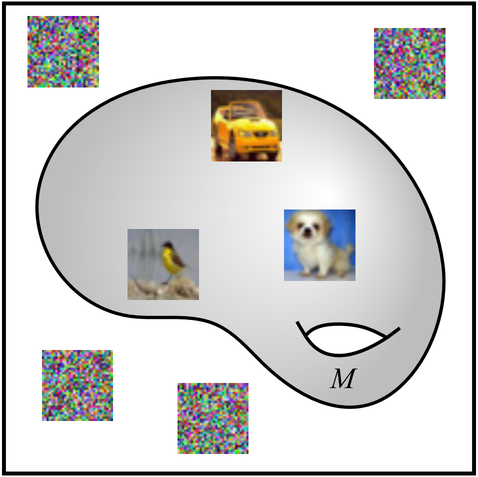

Manifold Learning

- Learn a model/representation for the data manifold
- Often involves finding a flat coordinate chart

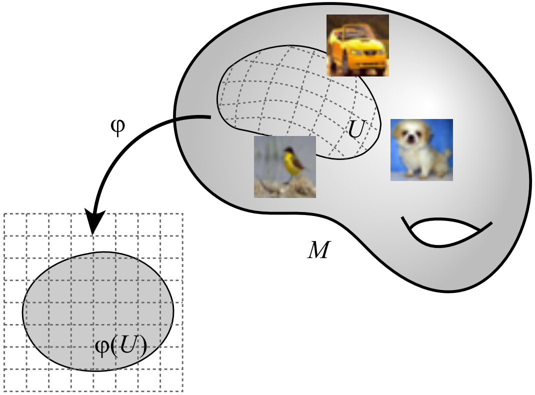

Manifold Learning

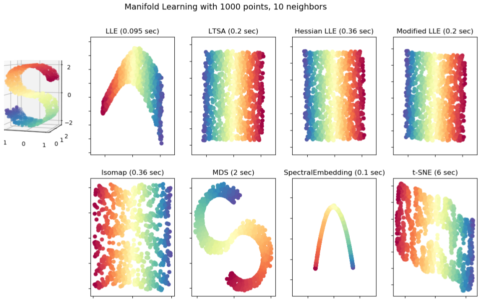

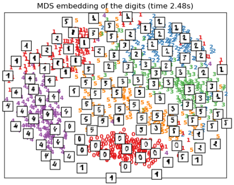

From scikit-learn.org

Volumes in High Dimensions

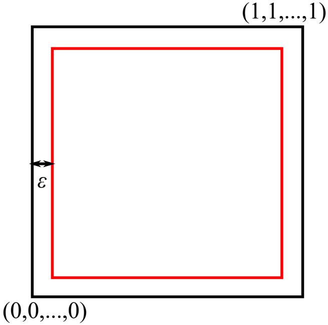

What is the volume of the unit d-cube shrunk by some small amount in each dimension?

$$V = (1 - 2\epsilon)^d$$

Approaches 0 as $d \to \infty$

Example: $256 \times 256 \times 3$ images, $\epsilon = \frac{1}{256}$

$$V \approx 2.0 \times 10^{-670}$$

Distances in High Dimensions

Sample two points uniformly from the unit d-cube:
$X, Y \sim \mathrm{Unif}([0,1]^d)$

What is the distribution of the distance between them?
$D = \|X - Y\|$

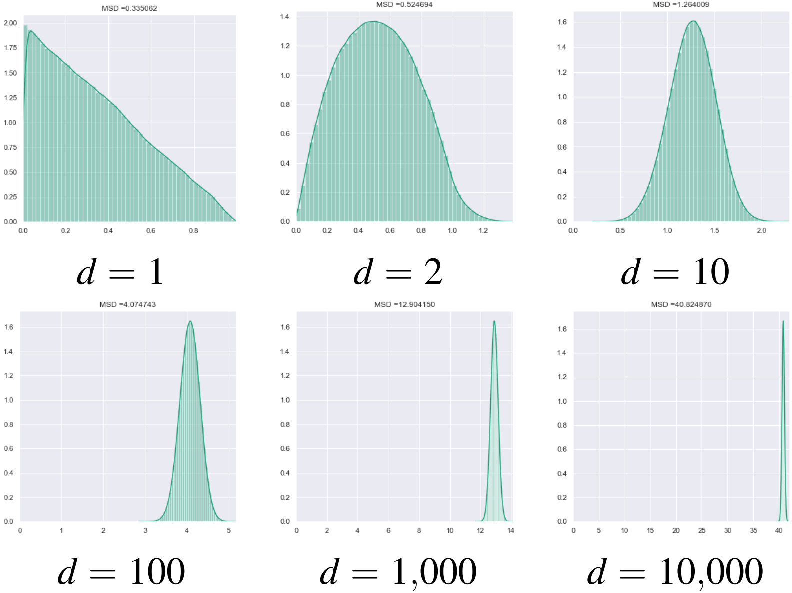
MSD =0.335062
$d = 1$
MSD =0.524694
$d = 2$
MSD =1.264009
$d = 10$
MSD =4.074743
$d = 100$
MSD =12.904150
$d = 1{,}000$
MSD =40.824870
$d = 10{,}000$

CIFAR-10

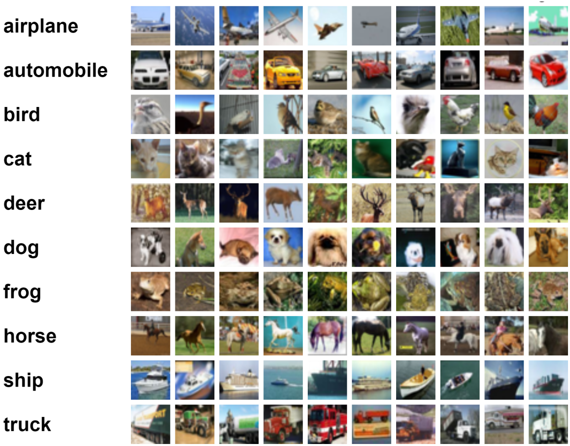

$32 \times 32 \times 3 = 3{,}072$ dimensions

10 classes

Distances in Real Data

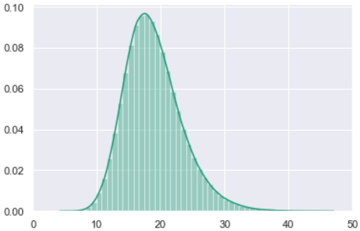

CIFAR-10

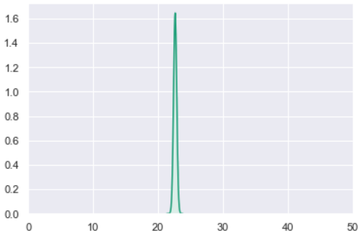

$\mathrm{Unif}([0, 1]^{3072})$

Manifold-valued Data

- Manifold already known, not learned
- Manifold arises from natural non-linear constraints on data
- Linear data analyses (in fact, vector space operations) violate these constraints

Directional Data

Data living on a circle (S^1) or sphere (S^2), etc.

- Orientation of molecules in protein structure
- Direction of robot or autonomous vehicle
- Position on the earth
- Motion capture: orientation of joints
- Time (time of day, day of the year, etc.)

Directional Data: Diffusion MRI

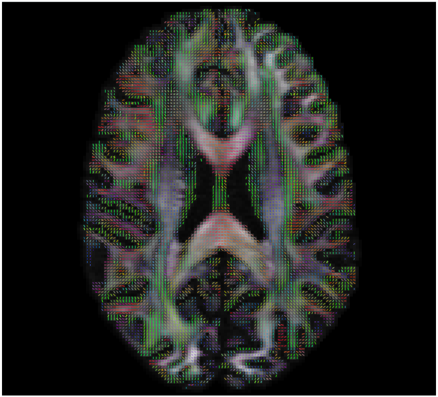

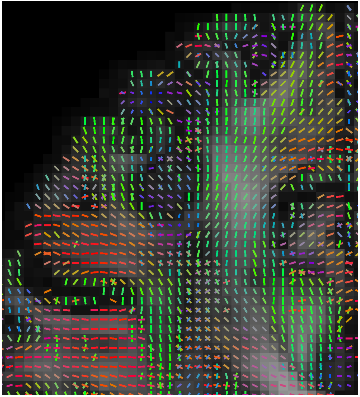

Voxel features are directions of axons in brain

Shape Manifolds

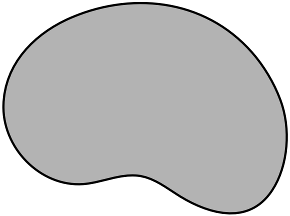

A shape is a point in a high-dimensional, nonlinear manifold, called a **shape space**.

Shape Manifolds

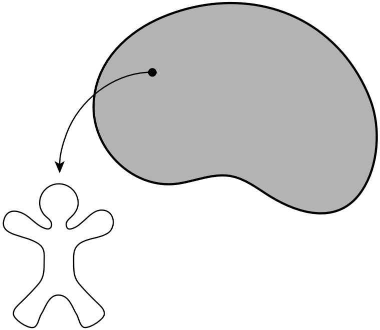

A shape is a point in a high-dimensional, nonlinear manifold, called a **shape space**.

Shape Manifolds

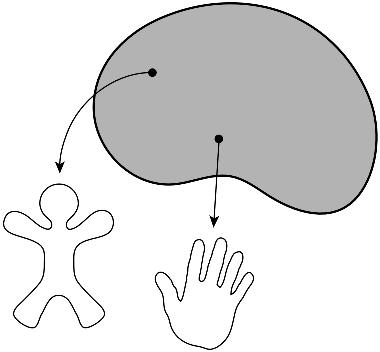

A shape is a point in a high-dimensional, nonlinear manifold, called a **shape space**.

Shape Manifolds

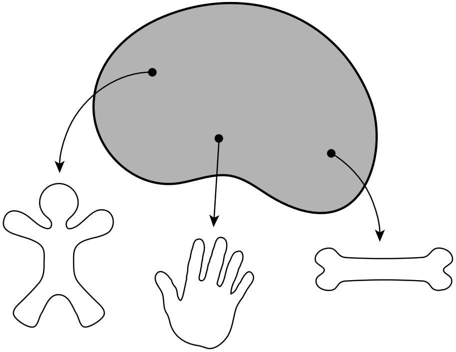

A shape is a point in a high-dimensional, nonlinear manifold, called a **shape space**.

Shape Manifolds

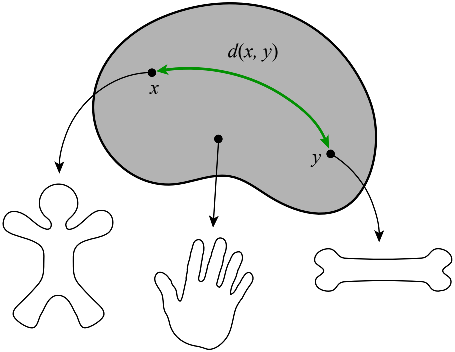

A metric space structure provides a comparison between two shapes.

Shape Statistics: Averages

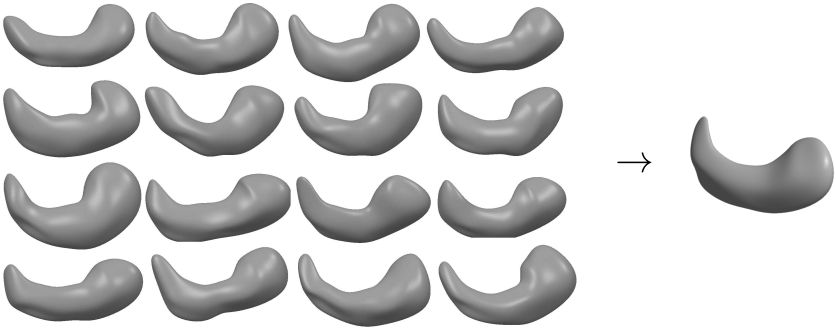

Shape Statistics: Variability

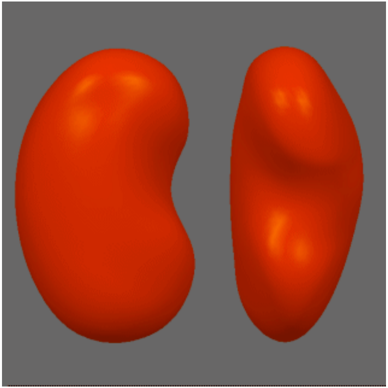

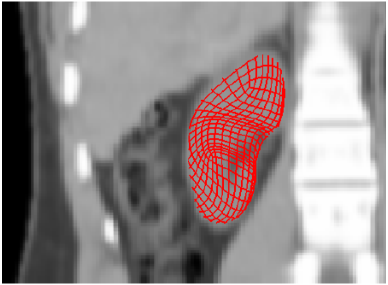

Shape priors in segmentation

Shape Application: Bird Identification

American Crow

Common Raven

Shape Statistics: Classification

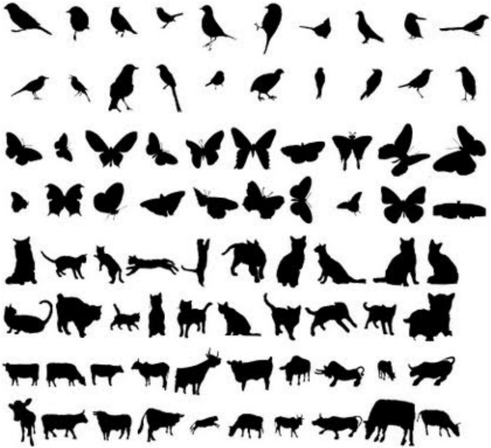

http://sites.google.com/site/xiangbai/animaldataset

Information Geometry

Parameters of a probability model live on manifolds

Example: covariance matrix of a 2D Gaussian distribution:

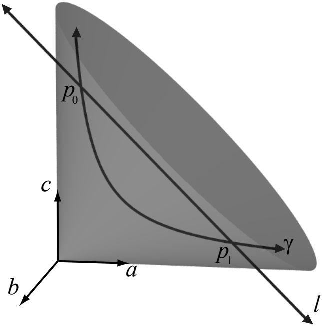

$\Sigma \in \mathrm{PD}(2)$ is of the form

$$\Sigma = \begin{pmatrix} a & b \\ b & c \end{pmatrix},$$

$$ac - b^2 > 0, \quad a > 0.$$

(positive-definite constraint)

Manifold Geometry of Deep Learning

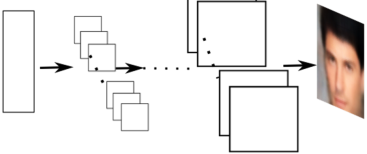

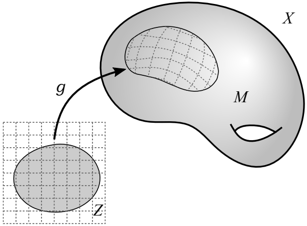

$$Z \subset \mathbb{R}^d \quad \xrightarrow{g} \quad X \subset \mathbb{R}^D$$

white noise → image